

Sciris: Simplifying scientific software in Python

Cliff C. Kerr^{1,2¶}, Paula Sanz-Leon¹, Romesh G. Abeysuriya^{1,3}, George L. Chadderdon^{3,4}, Vlad-Ştefan Harbuz⁵, Parham Saidi⁵, Maria del Mar Quiroga^{3,6}, Rowan Martin-Hughes³, Sherrie L. Kelly³, Jamie A. Cohen¹, Robyn M. Stuart^{1,7}, and Anna Nachesa⁸

¹ Institute for Disease Modeling, Global Health Division, Bill & Melinda Gates Foundation, Seattle, USA
² School of Physics, University of Sydney, Sydney, Australia ³ Burnet Institute, Melbourne, Australia ⁴ CAE USA, Tampa, USA ⁵ Saffron Software, Bucharest, Romania ⁶ Melbourne Data Analytics Platform, The University of Melbourne, Melbourne, Australia ⁷ Department of Mathematical Sciences, University of Copenhagen, Copenhagen, Denmark ⁸ Google, Zürich, Switzerland ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Sciris aims to streamline the development of scientific software by making it easier to perform common tasks. Sciris provides classes and functions that simplify access to frequently used low-level functionality in the core libraries of the scientific Python ecosystem (such as numpy for math and matplotlib for plotting), as well as in libraries of broader scope (such as multiprocessing for parallelization and pickle for saving and loading objects). While low-level functionality is valuable for developing robust software applications, it can divert focus from the scientific problems being solved. Some of Sciris' key features include: ensuring consistent dictionary, list, and array types (e.g., enabling users to provide inputs as either lists or arrays); enabling ordered dictionary elements to be referenced by index; simplifying datetime arithmetic by allowing date input in multiple formats, including strings; simplifying the saving and loading of files and complex objects; and simplifying the parallel execution of code. Sciris makes writing scientific code in Python faster, more pleasant, and more accessible, especially for people without extensive training in software development. With Sciris, users can achieve the same functionality with fewer lines of code, avoid reinventing the wheel, and spend less time looking up recipes on Stack Overflow. Sciris also forms the basis of ScirisWeb, an additional set of tools for building Flask-based Python webapps.

Statement of need

The landscape of scientific software

With the increasing availability of large volumes of data and computing resources, scientists across multiple fields of research have been able to tackle increasingly complex problems. But to harness these resources, the need for domain-specific software has become much greater. As the complexity of the questions being tackled has increased, so too has the software used to answer them, creating a steep learning curve and an increasing burden of code review ([Editorial Nature Methods, 2018](#)).

Scientific code workflows (e.g., either a full cycle in the development of a new software library, or in the execution of a one-off analysis) typically rely on multiple codebases, including but not limited to: low-level libraries, domain-specific open-source software, and self-developed and/or inherited Swiss-Army-knife toolboxes (whose original developer may or may not be around to pass on undocumented wisdom). Several scientific communities have adopted collaborative, community-driven, open-source software approaches due to the significant

savings in development costs and increases in code quality that they afford, such as [astropy](#) (Robitaille et al., 2013), [fmrip](#) (Esteban et al., 2019), and [nextstrain](#) (Hadfield et al., 2018). Despite this progress, a large fraction of scientific software development efforts remain a solo adventure (Kerr, 2019). This leads to proliferation of tools where resources are largely spent reinventing wheels of variable quality, which jeopardizes the code’s minimum requirements of being “re-runnable, repeatable, reproducible, reusable, and replicable” (Benureau & Rougier, 2018).

Beyond these requirements, low-level programming abstractions may get in the way of clarifying the science. For instance, one of the reasons PyTorch has become popular in academic and research environments is its success in making models easier to write compared to TensorFlow (Lorica, 2017). The need for libraries that provide “simplifying interfaces” for research applications is reflected in the development of multiple libraries in scientific Python ecosystems that have enabled researchers to focus their time and efforts on solving problems, prototyping solutions, deploying applications, and educating their communities. In addition to PyTorch (simplifying/extending Tensorflow), other examples include seaborn (simplifying/extending Matplotlib) (Waskom, 2021), pingouin (simplifying/extending pandas), and PyVista (simplifying/extending VTK) (Sullivan & Kaszynski, 2019), among many others.

Sciris in practice

The name [Sciris](#) is a portmanteau of “scientific” and “iris” (a reference to seeing clearly, as well as the Greek word for “rainbow”). We began work on it in 2014, initially to support development of the Optima suite of models Kerr et al. (2020). We repeatedly encountered the same inconveniences while building scientific webapps, and so we began collecting the tools we used to overcome them into a shared library. While Python is considered an easy-to-use language for beginners, the motivation that shaped Sciris’ evolution was to further lower the barriers to accessing the numerous supporting libraries we were using.

Our investments in Sciris paid off when in early 2020 its combination of brevity and simplicity proved crucial in enabling the rapid development of the Covasim model of COVID-19 transmission (Kerr et al., 2021). Covasim’s relative simplicity and readability, based in large part on its heavy use of Sciris, enabled it to become one of the most widely adopted models of COVID-19, used by students, researchers, and policymakers in over 30 countries (Kerr et al., 2022).

In addition to Covasim, Sciris is currently used by a number of other scientific software tools, such as [Optima HIV](#) (Kerr et al., 2015), [Optima Nutrition](#) (Pearson et al., 2018), the [Cascade Analysis Tool](#) (Kedziora et al., 2019), [Atomica](#) (The Atomica Team, 2020), [Optima TB](#) (Goscé et al., 2021), the [Health Interventions Prioritization Tool](#) (Fraser-Hurt et al., 2021), [SynthPops](#) (Mistry et al., in preparation), and [FPSim](#) (O’Brien et al., 2022).

Thus, we believe use of Sciris will result in more efficient scientific code production for solo developers and teams alike, including increased longevity of new scientific libraries (Perkel, 2020). Some of the key functional aspects that Sciris provides are: (i) brevity through simple interfaces; (ii) “dejargonification”; (iii) fine-grained exception handling; and (iv) version management. We expand on each of these below, but first provide a vignette that illustrates many of Sciris’ features.

Vignette

Compared with a domain-specific language like MATLAB, even relatively simple scientific code in Python can require significant “boilerplate”. This extra code can obscure the key logic of the scientific question being addressed.

For example, imagine that we wish to sample random numbers from a user-defined function with varying noise levels, save the intermediate calculations, and plot the results. In vanilla

Python, each of these operations is somewhat cumbersome. Figure 1 presents two functionally identical scripts; the one written with Sciris is considerably more readable and succinct.

This vignette illustrates many of Sciris' most-used features, including timing, parallelization, high-performance containers, file saving and loading, and plotting. For the lines of the script that differ, Sciris reduces the number of lines of code required from 33 to 7, a 79% decrease.

1	# Define random wave generator	1	# Define random wave generator
2	import numpy as np	2	import numpy as np
3		3	
4	def randwave(std, xmin=0, xmax=10, npts=50):	4	def randwave(std, xmin=0, xmax=10, npts=50):
5	np.random.seed(int(100*std)) # Ensure differences between runs	5	np.random.seed(int(100*std)) # Ensure differences between runs
6	a = np.cos(np.linspace(xmin, xmax, npts))	6	a = np.cos(np.linspace(xmin, xmax, npts))
7	b = np.random.randn(npts)	7	b = np.random.randn(npts)
8	return a + b*std	8	return a + b*std
9		9	
10	# Other imports	10	# Other imports
11	- import time	11	+ import sciris as sc
12	- import multiprocessing as mp		
13	- import pickle		
14	- import gzip		
15	- import matplotlib.pyplot as plt		
16	- from mpl_toolkits.mplot3d import Axes3D # Unused but must be imported		
17		12	
18	# Start timing	13	# Start timing
19	- start = time.time()	14	+ T = sc.timer()
20		15	
21	# Calculate output in parallel	16	# Calculate output in parallel
22	- multipool = mp.Pool(processes=mp.cpu_count())	17	+ waves = sc.parallelize(randwave, np.linspace(0, 1, 11))
23	- waves = multipool.map(randwave, np.linspace(0, 1, 11))		
24	- multipool.close()		
25	- multipool.join()		
26		18	
27	# Save to files	19	# Save to files
28	- filenames = []	20	+ filenames = [sc.save(f'wave{i}.obj', wave) for i, wave in enumerate(waves)]
29	- for i, wave in enumerate(waves):		
30	- filename = f'wave{i}.obj'		
31	- with gzip.GzipFile(filename, 'wb') as fileobj:		
32	- fileobj.write(pickle.dumps(wave))		
33	- filenames.append(filename)		
34		21	
35	# Create dict from files	22	# Create dict from files
36	- data_dict = {}	23	+ data = sc.odict({fname:sc.load(fname) for fname in filenames})
37	- for fname in filenames:		
38	- with gzip.GzipFile(fname) as fileobj:		
39	- filestring = fileobj.read()		
40	- data_dict[fname] = pickle.loads(filestring)		
41		24	
42	# Create 3D plot	25	# Create 3D plot
43	- data = np.array([data_dict[fname] for fname in filenames])	26	+ sc.surf3d(data[:,], plotkwargs=dict(cmap='orangeblue'))
44	- fig = plt.figure()		
45	- ax = plt.axes(projection='3d')		
46	- ny, nx = np.array(data).shape		
47	- x = np.arange(nx)		
48	- y = np.arange(ny)		
49	- X, Y = np.meshgrid(x, y)		
50	- surf = ax.plot_surface(X, Y, data, cmap='coolwarm')		
51	- fig.colorbar(surf)		
52		27	
53	# Print elapsed time	28	# Print elapsed time
54	- elapsed = time.time() - start	29	+ T.toc()
55	- print(f'Elapsed time: {elapsed:0.1f} s')		

Figure 1: Comparison of functionally identical scripts without Sciris (left) and with Sciris (right), showing a nearly five-fold reduction in lines of code. The resulting plots are shown in Figure 2.



Figure 2: Output of the scripts shown in Figure 1, without Sciris (left) and with Sciris (right). The two plots are identical except for the new high-contrast colormap available in Sciris.

Design philosophy

The aim of Sciris is to make common tasks simpler. Sciris includes implementations of heavily used code patterns and abstractions that facilitate the development and deployment of complex domain-specific scientific applications, and helps non-specialist audiences interact with these applications. We note that Sciris “stands on the shoulders of giants”, and as such is not intended as a replacement of these libraries, but rather as an interface that facilitates a more effective and sustainable development process through the following principles:

Brevity through simple interfaces. Sciris packages common patterns requiring multiple lines of code into single, simple functions. With these functions one can succinctly express and execute frequent plotting tasks (e.g., `sc.commaticks`, `sc.dateformatter`, `sc.plot3d`); ensure consistent types, including containers (e.g., `sc.toarray`, `sc.mergedicts`, `sc.mergelists`), or even perform line-by-line performance profiling (`sc.profile`). Brevity is also achieved by extending functionality of well established objects (e.g., `OrderedDict` via `sc.odict`) or methods (e.g., `isinstance` via `sc.checktype` that enables the comparison of objects against higher-level types like `arraylike`).

Dejargonification. Sciris aims to use plain function names (e.g., `sc.smooth`, `sc.findnearest`, `sc.safedivide`) so that the resulting code is as scientifically clear and human-readable as possible. Sciris also provides some **MATLAB**-like functionality, and uses the same names (e.g., `sc.tic` and `sc.toc`; `sc.boxoff`) to minimize the learning curve for scientists who have MATLAB experience.

Fine-grained exception handling. Across many classes and functions, Sciris uses the keyword `die`, enabling users to set a locally scoped level of strictness in the handling of exceptions. If `die=False`, Sciris is more forgiving and softly handles exceptions by using its default (opinionated) behavior, such as printing a warning and returning `None` so users can decide how to proceed. If `die=True`, it directly raises the corresponding exception and message. This flexibility reduces the need for try-catch blocks, which can distract from the code’s scientific logic.

Version management. Keeping track of dates, authors, and code versions, plus additional notes or comments, is an essential part of scientific projects. Sciris provides methods to easily save and load metadata to/from figure files, including Git information (`sc.savefig`, `sc.gitinfo`, `sc.loadmetadata`), as well as shortcuts for comparing module versions (`sc.compareversions`) or requiring them (`sc.require`).

Examples of key features

Here we illustrate a smattering of key features in greater detail; further information on installation and usage can be found at docs.sciris.org. Figure 3 illustrates the functional modules of Sciris. Sciris is available on pip (`pip install sciris`).

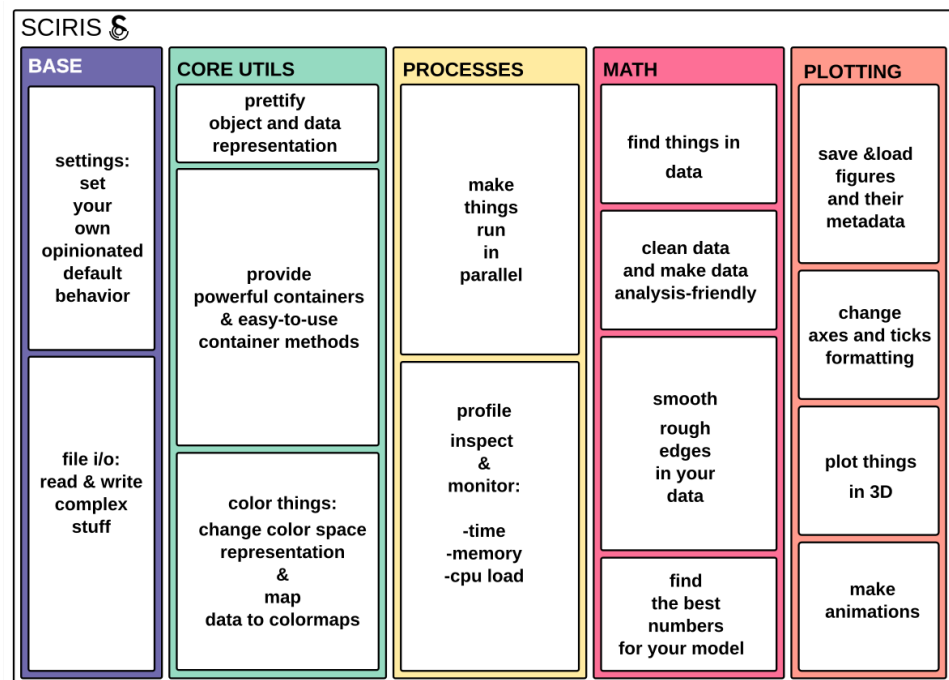


Figure 3: Block diagram of Sciris' functionality, grouped by high-level concepts and types of tasks that are commonly performed in scientific code.

High-performance containers

One of the key features in Sciris is `sc.odict`, a flexible container representing an associative array with the best-of-all-worlds features of lists, dictionaries, and numeric arrays. This is based on `OrderedDict` from [collections](https://pypi.org/project/collections), but supports list methods like integer indexing, key slicing, and item insertion:

```
> data = sc.odict(a=[1,2,3], b=[4,5,6])
> data['a'] == data[0]
> data[:].sum() == 21
> for i, key, value in data.enumitems():
    print(f'Item {i} is named "{key}" and has value {value}')
Item 0 is named "a" and has value [1, 2, 3]
Item 1 is named "b" and has value [4, 5, 6]
```

Numerical utilities

Indexing arrays is a common task in NumPy, but can be difficult due to incompatibilities of object type. `sc.findinds` will find matches even if two things are not exactly equal due to differences in type (e.g., floats vs. integers, lists vs. arrays). The code shown below produces the same result as calling `np.nonzero(np.isclose(arr, val))[0]`.

```
> sc.findinds([2,3,6,3], 3.0)
array([1,3])
```

140 Parallelization

141 A frequent hurdle scientists face is parallelization. Sciris provides `sc.parallelize`, which acts
142 as a shortcut for using `multiprocess.Pool()`. By default it adjusts the pool size based on the
143 CPUs available, but can also use either a fixed number of CPUs or allocate them dynamically
144 based on load (`sc.loadbalancer`). This example shows three equivalent ways to iterate over
145 multiple complex arguments:

```
> def f(x, y):
>     return x*y

> out1 = sc.parallelize(func=f, iterarg=[(1,2),(2,3),(3,4)])
> out2 = sc.parallelize(func=f, iterkwargs={'x':[1,2,3], 'y':[2,3,4]})
> out3 = sc.parallelize(func=f, iterkwargs=[{'x':1, 'y':2},
                                             {'x':2, 'y':3},
                                             {'x':3, 'y':4}])
```

146 Plotting

147 Numerous shortcuts for customizing and prettifying plots are available in Sciris. Several
148 commonly used features are illustrated below, with the results shown in [Figure 4](#):

```
> sc.options(font='Garamond') # Set custom font
> x = sc.daterange('2022-06-01', '2022-12-31', as_date=True) # Create dates
> y = sc.smooth(np.random.randn(len(x))*2)*1000 # Create smoothed random numbers
> c = sc.vectocolor(y, cmap='turbo') # Set colors proportional to y values
> plt.scatter(x, y, c=c) # Plot the data
> sc.dateformatter() # Automatic x-axis date formatter
> sc.commaticks() # Add commas to y-axis tick labels
> sc.setylim() # Automatically set the y-axis limits
> sc.boxoff() # Remove the top and right axis spines
```




Figure 4: Example of plot customizations via Sciris, including x- and y-axis tick labels and the font.

149 ScirisWeb

150 While a full description of [ScirisWeb](#) is beyond the scope of this paper, briefly, it builds on
151 Sciris to enable the rapid development of Python-based webapps, including those powering

152 [Covasim](#) and [Optima Nutrition](#). By default, ScirisWeb uses [Vuejs](#) and [sciris-js](#) for the frontend,
 153 [Flask](#) as the web framework, [Redis](#) for the (optional) database, and [Matplotlib/mpld3](#) for
 154 plotting. However, ScirisWeb is completely modular, which means that it could also be used to
 155 (for example) link a [React](#) frontend to a [MySQL](#) database with [Plotly](#) figures. This modularity
 156 is in contrast to full-stack solutions such as [Plotly Dash](#), [Streamlit](#), and [Voilà](#). While these
 157 libraries are even easier to use than ScirisWeb (since they do not require any knowledge of
 158 JavaScript), they provide limited options for customization or switching between technology
 159 stacks. In contrast, ScirisWeb provides the flexibility of a custom-written webapp within the
 160 context of an “it just works” framework.

161 Acknowledgements

162 The Sciris Development Team (info@sciris.org) wishes to thank David J. Kedziora, Dominic
 163 Delpont, Kevin M. Jablonka, Meikang Wu, and Dina Mistry for providing helpful feedback on
 164 the Sciris library. David P. Wilson, William B. Lytton, and Daniel J. Klein provided in-kind
 165 support of Sciris development. Financial and personnel support has been provided by the
 166 United States Defense Advanced Research Projects Agency (DARPA) Contract N66001-10-C-
 167 2008 (2010–2014), World Bank Assignment 1045478 (2011–2015), the Australian National
 168 Health and Medical Research Council (NHMRC) Project Grant APP1086540 (2015–2017), the
 169 Australian Research Council (ARC) Discovery Early Career Research Award (DECRA) Fellowship
 170 Grant DE140101375 (2014–2019), and the Institute for Disease Modeling (2019–present).

171 References

- 172 Benureau, F. C., & Rougier, N. P. (2018). Re-run, repeat, reproduce, reuse, replicate:
 173 Transforming code into scientific contributions. *Frontiers in Neuroinformatics*, 11, 69.
- 174 Editorial Nature Methods. (2018). Easing the burden of code review. *Nature Methods*, 15(9),
 175 641.
- 176 Esteban, O., Markiewicz, C. J., Blair, R. W., Moodie, C. A., Isik, A. I., Erramuzpe, A., Kent,
 177 J. D., Goncalves, M., DuPre, E., Snyder, M., & others. (2019). fMRIprep: A robust
 178 preprocessing pipeline for functional MRI. *Nature Methods*, 16(1), 111–116.
- 179 Fraser-Hurt, N., Hou, X., Wilkinson, T., Duran, D., Abou Jaoude, G. J., Skordis, J., Chukwuma,
 180 A., Lao Pena, C., Tshivuila Matala, O. O., Gorgens, M., & others. (2021). Using allocative
 181 efficiency analysis to inform health benefits package design for progressing towards universal
 182 health coverage: Proof-of-concept studies in countries seeking decision support. *PLOS*
 183 *One*, 16(11), e0260247.
- 184 Goscé, L., Abou Jaoude, G. J., Kedziora, D. J., Benedikt, C., Hussain, A., Jarvis, S., Skrahina,
 185 A., Klimuk, D., Hurevich, H., Zhao, F., & others. (2021). Optima TB: A tool to help
 186 optimally allocate tuberculosis spending. *PLOS Computational Biology*, 17(9), e1009255.
- 187 Hadfield, J., Megill, C., Bell, S. M., Huddleston, J., Potter, B., Callender, C., Sagulenko, P.,
 188 Bedford, T., & Neher, R. A. (2018). Nextstrain: Real-time tracking of pathogen evolution.
 189 *Bioinformatics*, 34(23), 4121–4123.
- 190 Kedziora, D. J., Abeyesuriya, R., Kerr, C. C., Chadderdon, G. L., Harbuz, V.-Ş., Metzger, S.,
 191 Wilson, D. P., & Stuart, R. M. (2019). The cascade analysis tool: Software to analyze and
 192 optimize care cascades. *Gates Open Research*, 3.
- 193 Kerr, C. C. (2019). Is epidemiology ready for big software? *Pathogens and Disease*, 77(1),
 194 ftz006.
- 195 Kerr, C. C., Stuart, R. M., Gray, R. T., Shattock, A. J., Fraser-Hurt, N., Benedikt, C., Haacker,
 196 M., Berdnikov, M., Mahmood, A. M., Jaber, S. A., & others. (2015). Optima: A model

- for HIV epidemic analysis, program prioritization, and resource optimization. *Journal of Acquired Immune Deficiency Syndromes*, 69(3), 365–376.
- Kerr, C. C., Stuart, R. M., Kedziora, D. J., Brown, A., Abeyasuriya, R., Chadderdon, G. L., Nachesa, A., & Wilson, D. P. (2020). Optima HIV methodology and approach. In F. Zhao, C. Benedikt, & D. Wilson (Eds.), *Tackling the world's fastest-growing HIV epidemic* (p. 291). The World Bank.
- Kerr, C. C., Stuart, R. M., Mistry, D., Abeyasuriya, R. G., Cohen, J. A., George, L., Jastrzebski, M., Famulare, M., Wenger, E., & Klein, D. J. (2022). Python vs. The pandemic: A case study in high-stakes software development. *Proceedings of the 21st Python in Science Conference (SciPy 2022)*.
- Kerr, C. C., Stuart, R. M., Mistry, D., Abeyasuriya, R. G., Rosenfeld, K., Hart, G. R., Núñez, R. C., Cohen, J. A., Selvaraj, P., Hagedorn, B., & others. (2021). Covasim: An agent-based model of COVID-19 dynamics and interventions. *PLOS Computational Biology*, 17(7), e1009149.
- Lorica, B. (2017). *Why AI and machine learning researchers are beginning to embrace PyTorch*. oreilly.com/radar/podcast/why-ai-and-machine-learning-researchers-are-beginning-to-embrace-pytorch
- Mistry, D., Kerr, C. C., Abeyasuriya, R. G., Wu, M., Fisher, M., Thompson, A., Skrip, L., Cohen, J. A., & Klein, D. J. (in preparation). *SynthPops: A generative model of human contact networks*.
- O'Brien, M. L., Valente, A., Chabot-Couture, G., Proctor, J., Klein, D., Kerr, C., & Zimmermann, M. (2022). FPSim: An agent-based model of family planning for informed policy decision-making. *PAA 2022 Annual Meeting*.
- Pearson, R., Killedar, M., Petravic, J., Kakietek, J. J., Scott, N., Grantham, K. L., Stuart, R. M., Kedziora, D. J., Kerr, C. C., Skordis-Worrall, J., & others. (2018). Optima nutrition: An allocative efficiency tool to reduce childhood stunting by better targeting of nutrition-related interventions. *BMC Public Health*, 18(1), 1–12.
- Perkel, J. M. (2020). Challenge to scientists: Does your ten-year-old code still run? *Nature*, 584(7822), 656–659.
- Robitaille, T. P., Tollerud, E. J., Greenfield, P., Droettboom, M., Bray, E., Aldcroft, T., Davis, M., Ginsburg, A., Price-Whelan, A. M., Kerzendorf, W. E., & others. (2013). Astropy: A community python package for astronomy. *Astronomy & Astrophysics*, 558, A33.
- Sullivan, C., & Kaszynski, A. (2019). PyVista: 3D plotting and mesh analysis through a streamlined interface for the visualization toolkit (VTK). *Journal of Open Source Software*, 4(37), 1450.
- The Atomica Team. (2020). Atomica: A simulation engine for compartmental models. In *GitHub repository*. GitHub. <https://github.com/atomicateam/atomica>
- Waskom, M. L. (2021). Seaborn: Statistical data visualization. *Journal of Open Source Software*, 6(60), 3021.